

Spousal Insurance and the Amplification of Business Cycles

Bence Bardóczy*

November 12, 2020

[Please click here for the most recent version.](#)

Abstract

I document that spousal labor supply substantially mitigates the impact of cyclical labor income risk on married households. Motivated by this evidence, I present a macroeconomic model with incomplete markets in which households are heterogeneous by gender and marital status. Couples can smooth their consumption over the business cycle better than singles because (i) spouses rarely lose their jobs at the same time; and (ii) secondary earners can increase their labor supply on the extensive margin in response to a job loss of the primary earner. According to my estimated model, spousal insurance mitigates the volatility of aggregate consumption by about 40%. Spousal insurance acts as a powerful automatic stabilizer because it weakens the general-equilibrium feedback between unemployment risk and economic activity. My model clarifies the circumstances under which this automatic stabilizer is stronger or weaker. Spousal insurance is particularly powerful in recessions caused by traditional demand shocks. It is less powerful in recessions caused by shocks like the current COVID epidemic.

1 Introduction

Households face large income uncertainty that varies systematically with the business cycle. This idiosyncratic risk is not fully insurable. [Blundell, Pistaferri and Saporta-Eksten \(2016\)](#) highlight three channels of partial insurance: (i) progressive taxes and transfers; (ii) accumulating assets; and (iii) risk sharing within the family. According to their estimates, the family channel is by far the most important form of insurance. In terms of consumption smoothing, it contributes more than the other two channels combined. Yet, macroeconomists have focused their efforts on the first two channels¹, whereas the business cycle implications of intra-family insurance have remained largely unexplored.

*Department of Economics, Northwestern University, 2001 Sheridan Road, Evanston, IL 60208. Email: bar-doczy@u.northwestern.edu. I am indebted to Martin Eichenbaum, Matthias Doepke, and Matthew Rognlie for their continuous guidance and support. I am very grateful to Adrien Auclert, Fabio Blasutto, Ana Danieli, João Guerreiro, Egor Kozlov, João Monteiro, Jane Olmstead-Rumsey, Christina Patterson, and Tímea Virágh for numerous helpful conversations. I thank participants at the Northwestern Macro Lunch for helpful comments. All errors are my own.

¹Examples for (i) include [Oh and Reis \(2012\)](#) and [McKay and Reis \(2016\)](#) who study the roles of targeted transfers and automatic fiscal stabilizers. Examples for (ii) include [Bayer, Lütticke, Pham-Dao and Tjaden \(2019\)](#) and [Kaplan, Moll and Violante \(2018\)](#) who focus on precautionary savings in liquid and illiquid assets.

In this paper, I argue that spousal insurance acts as an automatic stabilizer that dampens aggregate fluctuations. Recessions are amplified by countercyclical unemployment risk that encourages precautionary saving and depresses aggregate demand.² Spousal insurance dampens the propagation of unemployment risk in two ways. First, idiosyncratic risk is diversified within the family, which makes goods demand less responsive to unemployment. Second, secondary earners respond to the heightened risk of job loss to primary earners by increasing their labor supply. This precautionary labor supply—in sharp contrast to precautionary saving—reduces the cyclical volatility of output.

First, I discuss the empirical evidence. [Pruitt and Turner \(2020\)](#) provide direct evidence of spousal insurance from the joint tax returns of married couples. The administrative data show that household income is less volatile than individual income, and is less likely to fall by large amounts in recessions. They also confirm the observation that married women have less procyclical earnings than married men ([Doepke and Tertilt 2016](#)).

Using data from the Current Population Survey (CPS), I argue that differential patterns in the extensive margin of male and female employment can account for these facts. Three key features emerge from my empirical analysis. First, the probability of joint job loss by spouses is very low. Second, the probability of job loss is much more countercyclical for married men than married women. Third, labor force participation is procyclical for married men but countercyclical for married women.

Second, I build a dynamic stochastic general equilibrium (DSGE) model that is consistent with these facts and use it to characterize their implications for aggregate business cycles. The model has two key features. First, unemployment is determined in equilibrium in the presence of search frictions. Second, the extensive margin of spousal labor supply is a source of insurance against unemployment. In its other aspects, it is a standard incomplete markets model in the tradition of [Krusell and Smith \(1998\)](#).

The model economy is populated by three types of households: single men, single women, and married couples. I abstract from the dynamics of marriage and divorce and model these types as permanent. People of a particular type are ex-ante identical and face uninsurable labor income risk from two sources, idiosyncratic productivity shocks and job losses.

In each period, some jobs disappear exogenously. Workers who do not have a job decide whether to engage in costly job search or leave the labor force. Some workers who search do not receive an offer and are unemployed. Depending on their history, these people may or may not be eligible for unemployment benefits. People who leave the labor force continue to receive job offers, but at a lower rate than unemployed people.

Spousal insurance operates in the model via both a passive and an active channel. The passive channel reflects income pooling between employed spouses. The strength of this channel depends on the correlation of separation shocks between spouses. The active channel arises from family

²See [Ravn and Sterk \(2017\)](#) and [Challe, Matheron, Ragot and Rubio-Ramirez \(2017\)](#) for recent expositions of this argument that goes back to [Keynes \(1936\)](#).

labor supply decisions, which can take four different forms. First, an employed spouse can decide not to quit their job. Second, an unemployed spouse can continue to search for a job rather than leaving the labor force. Third, a spouse can start to search actively to increase their chances of finding a job. Fourth, a spouse who was not looking for a job may receive and accept a job offer.

According to the model, married households make joint labor supply choices that dampen the cyclical volatility of their labor income. The key here is that the labor force participation rate of secondary earners is countercyclical. In contrast, the labor force participation rate of single people is procyclical. These model-based outcomes are consistent with the evidence.

The intuition is the following. Secondary earners always have an incentive to enter the labor force when the primary earner is without a job. In a recession, more primary earners lose their jobs and are likely to be unemployed for a longer amount of time. As a result, non-employed secondary earners are more likely to enter the labor force during a recession. In contrast, non-employed singles are more likely to leave the labor force during a recession. The reason is that the job-finding probability is low. If they have enough savings, it is optimal for them to leave the labor force and re-enter when they receive an offer or when job-finding rates are back to normal.

Third, I use the quantitative model to isolate the role of spousal insurance in the propagation of aggregate shocks. I do so by considering a counterfactual economy in which married people do not provide insurance to each other. I subject both economies to the same aggregate demand shocks³ and compare the aggregate fluctuations they generate.

Spousal insurance implies less procyclical demand for goods and less procyclical supply of labor. The net result is a 42% decrease in the cyclical volatility of output. This effect is large for two reasons. First, married households account for a large share of aggregate consumption and labor supply. Second, there is a multiplier-like effect because unemployment risk is a function of output in the model.

The strength of the dampening effect depends critically on the cyclical volatility of job loss and job-finding rates. I estimate these to match the cyclical worker flows between labor market states using an exactly-identified Simulated Method of Moments (SMM). This is a challenging problem, because it requires simulating aggregate fluctuations in a complex heterogeneous-agent DSGE model. I develop a novel and efficient way of implementing SMM estimation in this setting.

Modeling spousal insurance explicitly is necessary for clarifying the circumstances under which it is stronger or weaker. First, the marriage rate among the working-age population determines the availability of spousal insurance. The decline in marriage observed in the United States since the 1960s, all else equal, has increased aggregate volatility. Second, the gender wage gap and the level of female labor force participation affect married women's ability to provide spousal insurance. In my model, a rise in the labor force participation of married women has opposing effects on passive and active insurance. As a result, I find that aggregate volatility is not sensitive to female labor force participation.

³I model these as shocks to the discount factor of households. Hall (2017) uses the same strategy to generate business cycles in a real model with labor search frictions.

Finally, I analyze the effectiveness of spousal insurance during the ongoing COVID-19 recession. In most recessions, fewer women lose their jobs than men. By contrast, as [Alon, Doepke, Olmstead-Rumsey and Tertilt \(2020a\)](#) show, in the COVID-19 recession more women lost their jobs than men. This unusual pattern also holds for married men and women. Moreover, the probability of both earners losing their jobs has been also unusually high during this recession. According to my model, this aspect of the COVID-19 shock rendered spousal insurance less effective than usual.

More generally, the example of the COVID-19 recession shows that the nature, and not just the size, of an economic shock matters for its impact on households. Aggregate shocks that originate in sectors with a high employment share of just one gender (such as construction) are better insured than shocks that affect women and men equally and are more likely to affect two spouses at the same time.

Literature. This paper relates and contributes to a diverse body of research on idiosyncratic income risk and its ramifications for households and the macroeconomy.

Improved computational power and access to administrative earnings data have facilitated the discovery of new facts about income inequality and risk ([Guvenen, Karahan, Ozkan and Song 2015](#)). The model I develop accommodates two recent findings and enables me to study their implications. First, [Hoffmann and Malacrino \(2019\)](#) document that large income shocks are associated with non-employment spells, which therefore drive the procyclical skewness of income growth first described by [Guvenen, Ozkan and Song \(2014\)](#). Second, [Pruitt and Turner \(2020\)](#) show that the joint income of married couples displays significantly less procyclical skewness than individual income, which points to the importance of family labor supply in mitigating cyclical income risk.⁴

Another line of research studies the various ways in which households insure themselves against income shocks ([Blundell, Pistaferri and Preston 2008](#), [Guler, Guvenen and Violante 2012](#), [Ortigueira and Siassi 2013](#), [Heathcote, Storesletten and Violante 2014](#)). [Blundell et al. \(2016\)](#) develop a structural econometric framework in which family labor supply, assets, and the tax system all have a role in accommodating income shocks. According to their estimates, family labor supply contributes more to consumption smoothing than the other two channels combined. My work is motivated by these household-level analyses. It complements them by arguing that family labor supply shapes aggregate business cycles, too.

The focus on aggregate consequences of spousal insurance is shared by a few recent papers. The most closely related is [Mankart and Oikonomou \(2017\)](#), who argue that countercyclical labor supply of secondary earners can explain the acyclicity of the aggregate labor force participation rate. At the household level, I model spousal insurance similarly. The crucial difference is

⁴[Busch, Domeij, Guvenen and Madera \(2020\)](#) report that “within-household smoothing does not seem effective at mitigating skewness fluctuations”. My conjecture is that they reach this conclusion because they exclude individuals with earnings below a minimum threshold. This selection criterion removes long unemployment spells and some movements on the participation margin. [Pruitt and Turner \(2020\)](#) keep individuals with zero earnings. My own analysis of the CPS data suggests that countercyclical labor force participation of secondary earners is an important channel of spousal insurance.

that, in my model, search frictions are determined endogenously, which creates a feedback between spousal insurance and unemployment risk. Moreover, I account for the profound gender differences in unemployment risk, rather than modeling individuals as ex-ante identical. Also emphasizing gender differences in labor market dynamics is [Ellieroth \(2019\)](#). She presents a partial equilibrium model with precautionary labor supply to rationalize married women’s countercyclical labor market flows. Last but not least, [Birinci \(2019\)](#) studies how alternative government transfers interact with spousal insurance over the business cycle.

In its objective and methodology, my paper belongs to the rapidly growing literature that studies how microeconomic heterogeneity affects the transmission of macroeconomic shocks ([McKay and Reis 2016](#), [Guerrieri and Lorenzoni 2017](#), [Kaplan et al. 2018](#), [Bayer et al. 2019](#), [Auclert, Rognlie and Straub 2018](#), [De Ferra, Mitman and Romei 2020](#)). This “HANK” literature is almost unified in its focus on income and wealth inequality between households.⁵ My results demonstrate that gender and marital status are key dimensions of household heterogeneity because they correlate with exposure to cyclical income risk. By modeling spousal insurance explicitly, I also take a step toward endogenizing the dynamics of income inequality.⁶ In terms of methodology, I extend a state-of-the-art solution method ([Auclert, Bardóczy, Rognlie and Straub 2019](#)) to the case of discrete-continuous choices, and work out a tractable way of estimating HANK models via Simulated Method of Moments. I hope these tools will be useful for other researchers in the field.

A key element of spousal insurance in my model is that secondary earners are more likely to enter the labor force when the primary earner loses their job. In labor economics, this is called the “added worker effect”, following the seminal work of [Lundberg \(1985\)](#). [Mankart and Oikonomou \(2016, 2017\)](#) document that the added worker effect has increased over time in the United States, and has been about 8% in the last two decades.⁷ Although the magnitude may seem low at first sight ([Gorbachev 2016](#), [Ellieroth 2019](#)), it is quite powerful through the lens of a dynamic model in which households are heterogeneous with respect to wealth. In my calibrated model, the added worker effect is just 5%, but entrants are selected from poor households who need extra income the most. Moreover, there is excess entry for months after the primary earner’s job loss, from households that run down their savings.

My work contributes to the literature on how long-term changes in women’s employment affect business cycle fluctuations ([Doepke and Tertilt 2016](#), [Fukui, Nakamura and Steinsson 2018](#)). [Albanesi \(2019\)](#) estimates that women’s rising share of aggregate hours and countercyclical la-

⁵A notable exception is [Patterson \(2018\)](#), who argues that business cycles are amplified by their unequal incidence with respect to demographic characteristics (gender, race, and age). She does not consider marital status or spousal labor supply.

⁶Although the distribution of income shocks plays a central role in this class of models, it is typically let to be determined by an exogenous Markov chain. Models with search and matching frictions ([Gornemann, Kuester and Nakajima 2016](#), [Den Haan, Rendahl and Riegler 2017](#), [Kekre 2019](#), [Graves 2019](#)) depart from this benchmark by endogenizing some aspect of unemployment (typically the job-finding rate). The richest model in this vein is [Alves \(2019\)](#)’s, which features on-the-job search and wage setting by firms who Bertrand compete for workers. State-dependent unemployment risk is a key feature of my model as well, where it is a main driver of spousal labor supply.

⁷A married woman whose husband loses his job in month t is 8 percentage points more likely to enter the labor force in month t than a woman whose husband remained employed.

bor supply played a crucial role in jobless recoveries, the productivity slowdown and the Great Moderation. My model accounts for the lower cyclical of women’s employment, partly as a result of an spousal insurance channel. It shows that higher female labor force participation does not necessarily reduce aggregate volatility, because it has opposing effects on passive and active insurance.

Finally, my paper makes contact with the literature on the macroeconomic consequences of the COVID-19 pandemic. I document that the epidemic led to job losses that were three times more highly correlated among spouses than usual. Together with the fact that women’s employment was hit especially hard, this implies that spousal insurance was weaker than in regular recessions. Here, my results complement [Alon, Doepke, Olmstead-Rumsey and Tertilt \(2020b\)](#), who emphasize that family labor supply was disrupted by the dramatic increase in childcare needs due to the widespread closure of schools and daycares.

2 Empirical Evidence

In this section, I present empirical evidence to show that spousal labor supply mitigates the labor income risk faced by married individuals. First, I review the latest evidence from administrative data on tax returns ([Pruitt and Turner 2020](#)). This data shows that (i) household income is less volatile than individual income; (ii) household income is less likely to fall by large amounts in recessions than individual income; (iii) married women have much less procyclical earnings than married men. Second, I use monthly worker flows to show that the extensive margin of employment can account for these facts. Specifically, I find that (a) women are systematically less likely to lose their job in typical downturns than men; (b) joint job loss by dual-earner couples is very rare; and (c) married women are less likely to leave the labor force during recessions. Collectively, these findings suggest that spousal insurance is effective against cyclical unemployment risk.

2.1 Earnings data

The best direct evidence for spousal insurance of labor income risk comes from the joint tax returns of married couples. My analysis is based on the aggregated dataset made available by [Pruitt and Turner \(2020\)](#).

Let y_{it} denote the gross labor earnings of individual i in year t , expressed in 2014 dollars. The variable I focus on is annual labor income growth, which is commonly used as an indicator of transitory labor income risk. This is computed as $x_{it} \equiv \log y_{it} - \log y_{it-1}$, where \$0 earnings are replaced with \$1. This measure is ideal for my purposes as it keeps individuals with occasionally zero earnings in the sample without leading to infinite growth rates. As we will see, movements on the extensive margin are important for both income risk and spousal insurance.

The two rows of figure 1 show two moments of income growth x_{it} : its standard deviation and

Kelley skewness⁸. The first two columns show these moments for the earnings of married men and women separately. The third column shows the moments for household income, which is simply the sum of the spouses' earnings. Each figure plots a given moment against the position of the household in the income distribution based on its average earnings in the full sample of 2000–2014. Finally, I report the statistics separately for recession years (2001–2002, 2008–2010) and expansion years (2003–2007, 2011–2014).

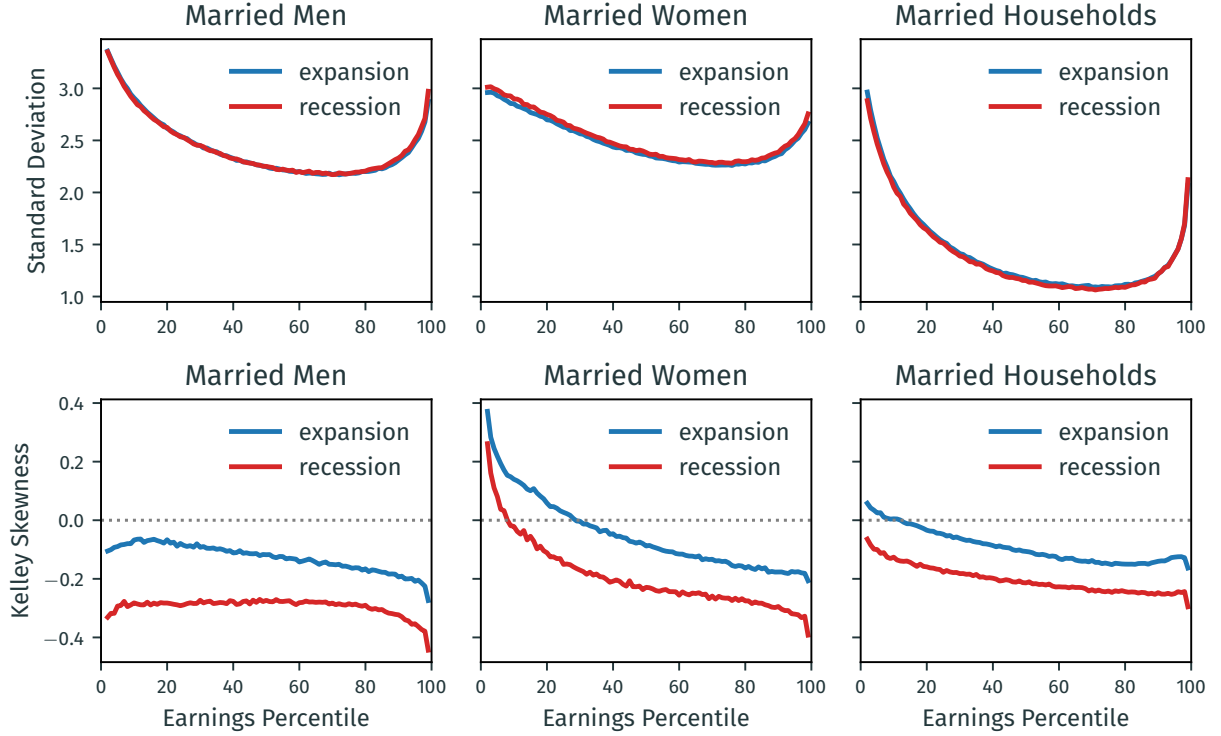


Figure 1: Distribution of One-Year Labor Income Growth for Individuals and Households

By comparing the graphs in the first line, we see that individuals face much more dispersion in their income growth than households. The standard deviation of a household's income growth is lower than that of its individual members across the income distribution, though the difference is largest for middle-income households. Furthermore, the standard deviation of income growth is remarkably acyclical.

This does not mean, however, that income risk itself is acyclical. Turning to skewness in the second line, we see that it is lower in recessions. That is, during recessions, large income losses become more likely and large gains become less likely. To put it simply, cyclical income risk

⁸Let p_n denote the n^{th} percentile of a random variable. Kelley skewness is defined as

$$\frac{(p_{90} - p_{50}) - (p_{50} - p_{10})}{p_{90} - p_{10}}.$$

This measure takes values in $[-1, 1]$ and is less sensitive to potential outliers than the fourth moment.

materializes in the probability of large shocks. This echoes the findings of [Guvenen et al. \(2014\)](#), who focused on men above a certain income threshold.

However, men and women are not affected equally in typical recessions. Skewness of men's income growth is more negative and falls more in recessions than women's. In low-income households, women's income growth actually displays positive skewness. These observations suggest that women are either less exposed to large income losses than men or make less procyclical labor supply choices. In the next section, I verify that both of these channels are in play by analyzing monthly worker flows in the Current Population Survey (CPS).⁹

2.2 Worker flows

Tax data paint an accurate picture of household income, but they do have some limitations. Most importantly, they provide little guidance to distinguish income shocks from labor supply choices. This is partly because the annual frequency of tax data is low relative to the frequency of income shocks and households' ability to respond to them. Take the example of job loss. In the United States, many employees can be dismissed at will or at most with one month's notice. In turn, the average unemployed person who is actively searching for a job can expect to find one in less than three months.

To address this challenge, I turn to the Current Population Survey (CPS). The CPS is a monthly labor force survey of about 60,000 households that is designed to be representative at the national level. Households are in the survey for 4 consecutive months, which allows me to construct monthly transition probabilities between employment (E), unemployment (U), and non-participation (N) disaggregated by gender and marital status. These worker flows provide additional insights into the incidence of cyclical income risk and the ways in which spousal insurance works. As such, they are also valuable to discipline the quantitative model I develop in section 3.

Figure 2 shows all six transition probabilities for 1976–2019, the longest span for which CPS micro data are publicly available. One of the most prominent features of these time series is the decline in the EN transitions by married women that lasted until the mid 1990s. Since my primary interest is business cycle fluctuations, I will base my discussion on the 1995–2019 period, by which time female labor force participation has plateaued. Although most of my observations apply to the entire sample, separating recent cycles from long-term trends makes the exercise cleaner.

Next, I condense the information in the time series of worker flows into just two numbers: their long-run average and cyclical. Motivated by [Doepke and Tertilt \(2016\)](#), I measure the latter as the unconditional elasticity with respect to detrended GDP. To be precise, I report β from the linear regression $\log p_t = \alpha + \beta \log Y_t + u_t$ where p_t is a quarterly average of a monthly transition rate, and Y_t is quarterly real GDP HP-filtered with a smoothing parameter of 1600.

⁹[Hoffmann and Malacrino \(2019\)](#) provide a strong reason for tracing these patterns to the extensive margin of employment. Using Italian social security data, they find that non-employment spells drive the fat tails of the earnings growth distribution for prime-age men, including essentially all of its procyclical skewness. I found the same for both men and women in the German employment history data (SIAB).

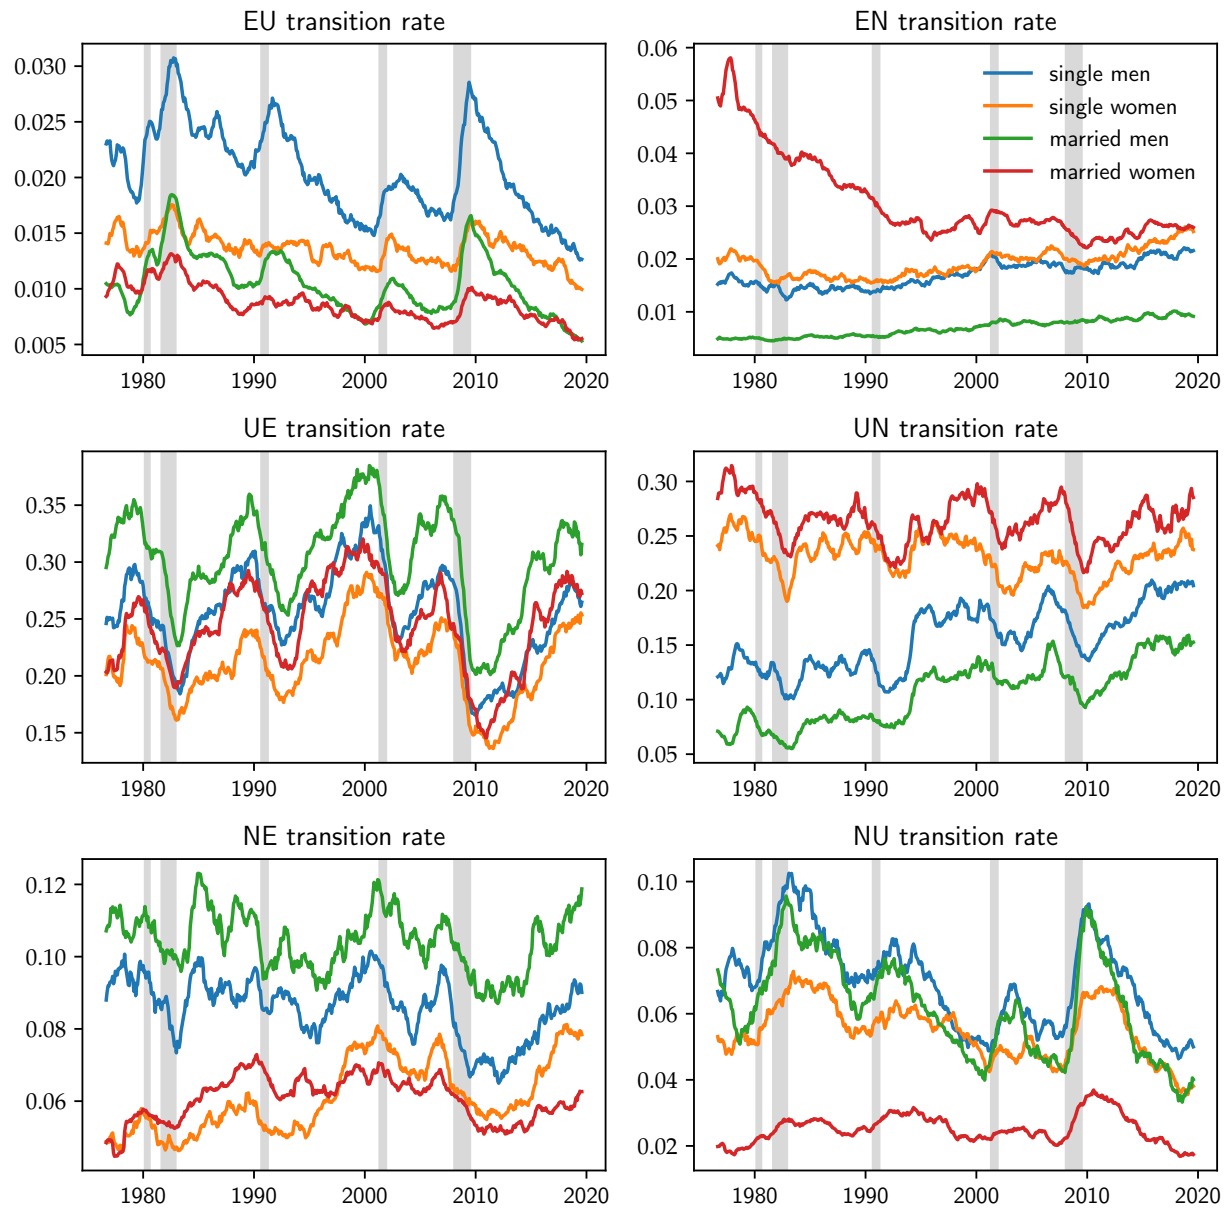


Figure 2: Worker flows by gender and marital status

Source: CPS monthly files. Sample is restricted to civilian population aged 25-54. Plotting 12-month centered moving averages of monthly transition probabilities.

Table 1: Outflows from Employment by Gender and Marital Status

		single men	single women	married men	married women
Average	<i>EU</i>	1.82%	1.29%	0.92%	0.77%
	<i>EN</i>	1.89%	2.06%	0.81%	2.59%
Cyclical	<i>EU</i>	-10.12*** (1.53)	-5.68*** (0.87)	-13.60*** (1.92)	-7.87*** (0.99)
	<i>EN</i>	0.90 (0.69)	1.28 (0.84)	-0.61 (0.89)	2.40** (0.74)

Source: Current Population Survey, monthly files for 1995–2019.

Sample: Civilian population aged 25–54.

Average: average of monthly transition rates.

Cyclical: elasticity to real GDP. HAC standard errors in parentheses. ***: p -value < 0.01 , **: p -value < 0.05 .

Table 1 shows information on the outflows from employment: *EU* and *EN*. The upper panel shows the average transition probabilities between 1995 and 2019. First, the magnitude of *EU* and *EN* flows are similar, which shows that they are both relevant. Second, the differences between the four groups are large: the average *EU* rate is more than twice as large for single men than for married women; and the average *EN* rate of married women is more than three times as large as that of married men.

The bottom panel of Table 1 shows that the *EU* rate is strongly countercyclical for all groups, with an almost twofold difference between men and women.¹⁰ In contrast, the *EN* rate is acyclical for singles and married men, and mildly procyclical for married women. This reflects a fundamental difference between the two outflows of employment. *EU* flows are mostly involuntary: resulting from layoffs, firings, and expiration of temporary contracts. Such events become more common in recessions. *EN* flows reflect a mix of choices (to not search actively) and acyclical shocks (e.g. to health). Notably, the estimated *EN* elasticity is highest and significant for married women. This indicates that delaying quitting in bad times is an active insurance channel that spouses provide.

Correlation of job loss within the family. If husbands lost their job every time their wives did, spousal insurance would be quite weak. Empirically, this is not the case. In the monthly CPS, I find that the correlation of *EU* transitions among dual-earner couples is just 0.042. Such a low value is consistent with the family economics literature. For example, exploiting the longer panel in the Survey of Income and Program Participation (SIPP), [Shore and Sinai \(2010\)](#) report that the probability that a married couple has overlapping unemployment spells in a year is 1.4% on average.¹¹ This shows that dual-earner couples who pool their income enjoy passive insurance against job loss simply because these shocks are largely uncorrelated within the family.

¹⁰[Albanesi and Şahin \(2018\)](#) estimate that gender difference in industry composition accounts for most of the difference in payroll employment changes during recessions.

¹¹See their table 3, column (4). They emphasize that the probability of joint unemployment is larger for couples who have the same occupation, but such couples are only 3% of their sample.

Table 2: Inflows to Employment by Gender and Marital Status

		single men	single women	married men	married women
Average	UE	25%	22%	30%	25%
	NE	8%	7%	10%	6%
Cyclical	UE	10.71*** (1.32)	8.72*** (1.55)	10.36*** (1.17)	8.40*** (1.55)
	NE	4.74*** (1.06)	3.74*** (1.00)	2.78** (0.91)	2.30* (0.75)

Source: Current Population Survey, monthly files for 1995–2019.

Sample: Civilian population aged 25–54.

Average: average of monthly transition rates.

Cyclical: elasticity to real GDP. HAC standard errors in parentheses. ***: p -value < 0.01, **: p -value < 0.05.

Table 2 shows an analogous summary of employment inflows: *UE* and *NE*. The upper panel shows two well-known facts about US labor markets. First, monthly *UE* flows are high, which translates into short unemployment durations. Second, *NE* rates are far from negligible. Given that there are many more non-participants than unemployed, the total number of *NE* flows is actually higher than of *UE* flows. This observation reinforces the importance of the participation margin for the determination of household income.

Turning to cyclicalities, we see that *UE* flows are strongly procyclical for all groups, and more similar than *EU* flows. Men have somewhat more procyclical job-finding rates than women, but the difference is not too large. *NE* rates are also procyclical, albeit to a much lesser degree. This makes sense. *NE* transitions require finding a job, which is harder in recessions. However, a large share of non-participants don't want a job, which mitigates the effect of the underlying job-finding rate to *NE* flows. In terms of spousal insurance, notice that married non-participants have significantly less procyclical *NE* rates, which is consistent with an added worker effect subject to countercyclical job-finding frictions (Lundberg 1985).

Table 3: Flows on the Participation Margin

		single men	single women	married men	married women
Average	<i>UN</i>	18%	23%	13%	26%
	<i>NU</i>	6%	5%	5%	2%
Cyclical	<i>UN</i>	5.91*** (1.00)	4.09*** (0.81)	6.64*** (1.22)	5.43*** (0.64)
	<i>NU</i>	-10.06*** (1.34)	-6.55*** (1.23)	-13.98*** (1.47)	-8.73*** (1.23)

Source: Current Population Survey, monthly files for 1995–2019.

Sample: Civilian population aged 25–54.

Average: average of monthly transition rates.

Cyclical: elasticity to real GDP. HAC standard errors in parentheses. ***: p -value < 0.01 , **: p -value < 0.05 .

Finally, table 3 shows the flows between unemployment and non-participation. Note that these flows are somewhat less reliable than the flows on the employment margin, because the CPS is prone to classification error between *U* and *N* (Abowd and Zellner 1985). With this caveat in mind, the *UN* flows are procyclical across the board, though slightly more so for the married than singles. Elsby, Hobijn and Şahin (2015) attribute this to a composition effect. In recessions, a large number of workers who lose their jobs have high attachment to the labor market and thus are less likely to leave the labor force.

NU flows in turn are countercyclical and significantly more so for married people. Again, this difference is consistent with the added worker effect: non-working spouses enter in bad times when the primary earner faces more risk. The conventional wisdom is that the added worker effect should be stronger for women, who are more likely to be secondary earners. This is not what we see in aggregate *NU* flows: they are more cyclical for married men than women and substantial for singles as well. This is sometimes interpreted as evidence that the added worker effect is weak (Elsby et al. 2015). This is certainly true in the sense that the added worker effect cannot be the dominant driver of aggregate *NU* flows. However, it does not mean that the added worker effect is weak at the household level. In the aggregate, it may be masked by the fact that, among non-participants, the average married woman is less attached to the labor force than the average married man. In section 5, I show that my calibrated model is consistent with the added worker effect at the household level, estimated by Mankart and Oikonomou (2017).

Takeaway. In recessions, job loss becomes more common and new jobs become harder to find which leads to longer unemployment durations. Countercyclical unemployment accounts for the procyclical skewness of annual earnings growth. Exposure to cyclical unemployment risk is not uniform: men are more exposed than women, and single people are more exposed than married people. Flows on the participation margin, which are more likely to reflect labor supply choices than shocks, are consistent with spousal insurance provision.

3 A business cycle model with spousal insurance

In this section, I introduce a macroeconomic model with heterogeneous households and incomplete markets, in which spousal labor supply is a source of insurance against idiosyncratic shocks.

The economy consists of a household sector, a firm sector and public sector. Households consume, supply labor to firms, and save in government bonds. They also face uninsurable income risk that includes unemployment. There is a representative final good producer who uses labor to produce a homogeneous final good. The fiscal authority collects labor income tax and issues one-period real bonds to finance its expenditures on the final good, lump-sum transfers and unemployment benefits.

3.1 Households

There is a unit mass of infinitely-lived households of three types: single men, single women, and married couples. For simplicity, I assume that gender and marital status are permanent, that is, there are neither new marriages nor divorce. Married couples make all their decisions jointly.

All households have time-separable preferences over consumption and leisure and face uninsurable labor income risk from unemployment and productivity shocks. The parameterization of the utility function and the income process may differ by gender and marital status, but households of the same type are all identical *ex ante*. Nevertheless, idiosyncratic shocks generate income and wealth inequality among them.

State variables. A household is characterized by the employment status of its members $s \in \{E, U_b, U_{nb}, N\}$ (employed, unemployed with benefits, unemployed with no benefits, non-participant), the labor productivity of its members $z \in \mathcal{G}$ (a finite set) and assets $a \geq \underline{a}$.

Employment status affects utility as well as income, because market work and job search take time away from leisure (or home production). Specifically, let the flow utility of an individual in employment state s_{it} who consumes c_{it} be given by

$$u(s_{it}, c_{it}) = \begin{cases} \log(c_{it}) - \varphi & \text{for } s_{it} = E, \\ \log(c_{it}) - \chi & \text{for } s_{it} = U_b, U_{nb}, \\ \log(c_{it}) & \text{for } s_{it} = N, \end{cases} \quad (1)$$

where $\varphi > \chi > 0$ are the utility costs of formal employment and active job search, respectively. Married couples maximize their joint utility, which is simply the sum of their individual utilities. Since utility from consumption is concave and separable from leisure, it is always optimal for both spouses to have the same consumption.

The post-tax labor income of an individual in employment state s_{it} with labor productivity z_{it} ,

given real wage w_t , labor tax τ_t , and transfers T_t , is given by

$$y_t(s_{it}, z_{it}) = \begin{cases} (1 - \tau_t)w_t z_{it} + T_t & \text{for } s_{it} = E, \\ (1 - \tau_t)w_t \min\{bz_{it}, \bar{b}\} + T_t & \text{for } s_{it} = U_b, \\ T_t & \text{for } s_{it} = U_{nb}, N. \end{cases} \quad (2)$$

Unemployment benefits are specified as a constant replacement rate b up to a cap \bar{b} . In addition to labor income, households also earn financial income of $r_t a_{it-1}$. Married couples pool their individual labor income and own assets jointly. That is, the labor income of a married couple is the sum of the husband's and the wife's income.

Evolution of state variables. Household i enters period t with state variables inherited from last period $(s_{it-1}, z_{it-1}, a_{it-1})$.

1. Productivity shock: z_{it} is drawn according to an exogenous Markov chain Π .
2. Employment shocks: each member of the household moves to an interim employment state $x_{it} \in \{M, B, L\}$ (matched with a firm, unmatched but eligible for UI, unmatched and ineligible for UI). The transition probabilities are shown by table 4 below.

	matched (M)	UI eligible (B)	neither (L)
E	$1 - s_t$	s_t	0
U_b	f_t	$(1 - \zeta)(1 - f_t)$	$\zeta(1 - f_t)$
U_{nb}	f_t	0	$1 - f_t$
N	f_t^*	0	$1 - f_t^*$

Table 4: Transition matrix of employment shocks

An employed worker may lose her job with probability s_t and become UI eligible. Unemployed workers find a job with probability f_t . If they don't find a job, their unemployment benefits expire with probability ζ .¹² An unemployed without benefits has to cycle through an employment spell to qualify for UI again. Non-participants get matched with a lower probability $f_t^* < f_t$.

3. Labor supply choice: Discrete choice that is constrained by the outcome of the employment shock in the previous stage. Table 5 shows the choice set for each interim state x_{it} .

¹²This is a parsimonious way to capture the limited duration of unemployment benefits.

	E	U_b	U_{nb}	N
M	✓		✓	✓
B		✓	✓	✓
L			✓	✓

Table 5: Feasible labor supply choices

Only matched workers can become employed, and those who reject a job cannot collect UI.

4. Consumption-savings choice Income is received and split between consumption c_{it} and gross savings a_{it} . Household i finishes period t in the new state (s_{it}, z_{it}, a_{it}) .

The decision problem of singles. Let $W_{gt}(x_{it}, z_{it}, a_{it-1})$ denote time- t value function of a single household of gender g in stage 2, just after the exogenous productivity and employment shocks have been realized. Let $V_{gt}(s_{it}, z_{it}, a_{it-1})$ be the value function in stage 3, after the discrete choice has been made.

Armed with these definitions, we can write the Bellman equation in two steps. First,

$$W_{gt}(x_{it}, z_{it}, a_{it-1}) = \max_{s_{it} \in \Gamma(x_{it})} V_{gt}(s_{it}, z_{it}, a_{it-1}) \quad (3)$$

where $\Gamma(x_{it}) \subset \{E, U_b, U_{nb}, N\}$ is the set of feasible labor supply choices in state x_{it} given by table 5. Second,

$$V_{gt}(s_{it}, z_{it}, a_{it-1}) = \max_{c_{it}, a_{it}} \left\{ u_g(s_{it}, c_{it}) + \beta_{gt} \mathbb{E}_t [W_{gt+1}(x_{it+1}, z_{it+1}, a_{it})] \right\}, \quad (4)$$

s.t. $c_{it} + a_{it} = y_t(s_{it}, z_{it}) + (1 + r_t)a_{it-1},$
 $a_{it} \geq \underline{a}.$

Note that the gender-specific parameters—indexed by g —are the utility parameters $\{\beta_g, \varphi_g, \chi_g\}$, the labor market frictions $\{f_{gt}, f_{gt}^*, s_{gt}\}$, and the Markov process for labor productivity $\{\mathcal{G}_g, \Pi_g\}$.

The decision problem of couples. Since married households pool their incomes and assets to maximize their joint utility, the structure of their decision problem is analogous to singles', just with a larger state space to account for all combinations of individual labor market outcomes.

To be specific, couple i is characterized by a pair of employment states $s_{it} = (s_{it}^m, s_{it}^f) \in \{E, U_b, U_{nb}, N\}^2$, a pair of labor productivities $z_{it} = (z_{it}^m, z_{it}^f) \in \mathcal{G}^m \times \mathcal{G}^f$, shared assets $a_{it-1} \geq \underline{a}$.

The first part of the Bellman equation is

$$W_t(x_{it}, z_{it}, a_{it-1}) = \max_{s_{it} \in \Gamma(x_{it})} V_t(s_{it}, z_{it}, a_{it-1}) \quad (5)$$

where the choice sets $\Gamma(x_{it}) \subset \{E, U_b, U_{nb}, N\}^2$ are a straightforward generalization of table 5 to two household members. For example, if the husband is unmatched without UI and the wife is

matched, their choice set is $\Gamma(BM) = \{U_{nb}E, U_{nb}U_{nb}, U_{nb}N, NE, NU_{nb}, NN\}$. The second part is given by

$$V_t(s_{it}, z_{it}, a_{it-1}) = \max_{c_{it}, a_{it}} \left\{ u_m(s_{it}^m, c_{it}) + u_f(s_{it}^f, c_{it}) + \beta_t \mathbb{E}_t [W_{t+1}(x_{it+1}, z_{it+1}, a_{it})] \right\}, \quad (6)$$

$$\text{s.t. } c_{it} + a_{it} = y_t(s_{it}^m, z_{it}^m) + y_t(s_{it}^f, z_{it}^f) + (1 + r_t)a_{it-1},$$

$$a_{it} \geq \underline{a}.$$

Labor market transitions at the household level are obtained, for the most part, by multiplying the individual transition probabilities from table 4. For example, the probability that an U_bE couple becomes an MM couple next period is $f_{t+1}^m(1 - s_{t+1}^f)$ (husband finds a job and wife does not lose hers). That is, most employment shocks are uncorrelated between spouses. The exceptions are EE couples, for whom I allow for an arbitrary correlation ρ of job loss. In appendix C.1, I show how to construct the household-level transition rates from the average separation rates of married men and women.

3.2 Firms

My priority is to capture the observed magnitude of cyclical unemployment risk. This requires endogenizing separation rates (frequency of unemployment spells) as well as job-finding rates (duration of unemployment spells). As we saw in section 2.2, both margins are strongly cyclical and heterogeneous by gender and marital status. Therefore, I opt for a simple and flexible setup that allows me to specify these moments almost directly.

Final good firm. The final good market is competitive. A representative firm produces a homogeneous final good using a linear technology

$$Y_t = \Theta_t L_t. \quad (7)$$

The firm hires workers subject to search frictions, which it takes as given. The identity of the workers is irrelevant for the firm, who pays each worker their marginal product and makes zero profits.

Search frictions. Let hatted variables denote log deviations from steady state. The labor market transition probabilities in table 4 are specified directly as a function of contemporaneous output fluctuations with type-specific elasticities $\epsilon_j(\bullet)$:

$$\hat{f}_{jt} = \epsilon_j(f) \cdot \hat{Y}_t, \quad \hat{f}_{jt}^* = \epsilon_j(f^*) \cdot \hat{Y}_t, \quad \hat{s}_{jt} = \epsilon_j(s) \cdot \hat{Y}_t. \quad (8)$$

Although I offer no microfoundation for these equations, it is worth noting that any specification that generates the same, empirically realistic, transition probabilities would have the same distributional effect on households and prompt equivalent precautionary responses.

3.3 Government

The fiscal authority issues one-period bonds B_t and levies a proportional labor tax τ_t on employed workers to finance its expenditures. These include lump-sum transfers T_t , unemployment benefits, and public consumption of the final good G_t . The government budget is thus given by

$$\tau_t w_t \int_E z_{it} di + B_t = (1 + r_t) B_t + G_t + T_t (1 + \omega_{mc}) + \underbrace{(1 - \tau_t) w_t \int_{U_b} \min\{bz_{it}, \bar{b}\} di}_{\text{total UI claims}}, \quad (9)$$

where ω_{mc} is the share of married couples. There is a unit mass of households, but every couple has two members and hence the total measure of individuals is $1 + \omega_{mc}$.

I treat G_t and T_t as exogenous, and assume that the government adjusts the proportional tax τ_t according to the rule

$$\tau_t - \tau_{ss} = \phi \frac{B_{t-1} - B_{ss}}{Y_{ss}}. \quad (10)$$

That is, taxes are increased whenever debt is above its steady state level, and the parameter $\phi > 0$ determines the speed of fiscal adjustment. Assuming smooth fiscal policy is useful in non-Ricardian models such as this to avoid the unreasonably large income effects that would result from balancing the budget period by period.

3.4 Market clearing

Labor, asset, and goods market clearing are given by

$$L_t = \int_E z_{it} di, \quad (11)$$

$$B_t = \int a_{it} di, \quad (12)$$

$$Y_t = C_t + G_t. \quad (13)$$

4 Numerical implementation

The general equilibrium in the model from section 3 depends on the distribution of households over individual states, an infinite-dimensional object. In addition, the presence of discrete choices makes my model more challenging to compute than typical heterogeneous-agent DSGE models. Therefore, I offer a brief overview of my approach to solve and estimate it.

4.1 Household problem

The decision problems of singles (3)(4) and couples (5)(6) have analogous structures and can be solved independently, since gender and marital status are permanent types.

The main complication is that discrete labor supply choices lead to kinks in the value functions and jumps in the policy functions for consumption and assets. In general, the points of discontinuity would not be on the grid, and there would be large approximation errors. Therefore, I follow [Iskhakov, Jørgensen, Rust and Schjerning \(2017\)](#) and introduce additive choice-specific independent and identically distributed extreme value taste shocks. An arbitrarily small taste shock is sufficient to make the policy functions continuous and admit accurate grid-based solutions. In sum, the Bellman equations (3) and (5) are modified to

$$W_t(x_{it}, z_{it}, a_{it-1}) = \max_{s_{it} \in \Gamma(x_{it})} \left\{ V_t(s_{it}, z_{it}, a_{it-1}) + \varepsilon(s_{it}) \right\} \quad (14)$$

where $\varepsilon(s_{it})$ is the realization of the taste shock for choice s_{it} .

Small taste shocks eliminate jumps but preserve non-convexities. As a result, first order conditions are necessary but not sufficient and the original endogenous gridpoint method of [Carroll \(2006\)](#) has to be extended with an upper envelope step. [Druehl \(2020\)](#) provides a useful guide on how to implement this efficiently.

4.2 Aggregate fluctuations

When it comes to general equilibrium models with substantial complexity both at their micro- and macro level, our most powerful methods resort to certainty equivalence with respect to aggregate shocks to ensure they can accurately capture the effects of large idiosyncratic uncertainty.¹³

In this paper, I extend the sequence-space Jacobian method of [Auclert et al. \(2019\)](#) to the case of discrete-continuous dynamic programs with taste shocks. This method stands out in its transparency as it requires no dimension reduction to be applied to large-scale problems. [Boppart, Krusell and Mitman \(2018\)](#) promote the use of small unanticipated (“MIT”) shocks to characterize aggregate fluctuations for this reason. The sequence-space Jacobian method is a general and systematic way of computing impulse responses to infinitesimal aggregate shocks.

The collective decisions of households are represented by sequence-space Jacobians. These are 1000×1000 matrices that contain all information from the cross-sectional distribution (discretized over 78,000 gridpoints) that is relevant for aggregate dynamics. The size of the Jacobians is determined by the truncation horizon. That is, I account for the effects of anticipating aggregate shocks up to 1000 months ahead as well as for the delayed effects of past shocks that occurred less than 1000 months ago. Increasing the truncation horizon has no discernible impact on my results.

Moment matching. The sequence-space Jacobian method is particularly well-suited for calibration and estimation, too. As I explain in section 5.3, I calibrate search frictions to match the cyclical volatility of flows between employment and unemployment. Here, I explain how to construct these moments in the model with aggregate shocks without resorting to simulation.

Let dX_t denote the percentage point deviation of a generic variable X_t from its deterministic

¹³This approach was pioneered in state space by [Reiter \(2002, 2009\)](#) and extended recently by [Winberry \(2018\)](#), [Ahn, Kaplan, Moll, Winberry and Wolf \(2018\)](#), [Bayer and Lueticke \(2020\)](#).

steady state. The elasticity of any transition probability p_t with respect to output Y_t is, by definition,

$$\varepsilon(p) = \frac{\text{Cov}(\text{d}p_t, \text{d}Y_t)}{\text{Var}(\text{d}Y_t)} \quad (15)$$

In the model, any deviation from steady state is a result of discount factor shocks. That is, $\text{d}p_t$, say, can be written in MA(∞) form as

$$\text{d}p_t = \sum_{s=0}^{\infty} \underbrace{\left(\frac{\text{d}p_t}{\text{d}\varepsilon_{t-s}^{\beta}} \right)}_{m_s^{p,\beta}} \varepsilon_{t-s}^{\beta} \quad (16)$$

where $\varepsilon_{t-s}^{\beta}$ are the past realizations of the discount factor shock. It follows that the covariances in the elasticity formula (15) can be expressed as

$$\text{Cov}(\text{d}p_t, \text{d}Y_t) = \sigma_{\beta}^2 \sum_{s=0}^{\infty} m_s^{p,\beta} m_s^{Y,\beta}, \quad \text{and} \quad \text{Var}(\text{d}Y_t) = \sigma_{\beta}^2 \sum_{s=0}^{\infty} (m_s^{Y,\beta})^2. \quad (17)$$

Recall that I solve the model assuming certainty equivalence with respect to aggregate shocks. As a corollary, the MA coefficients $m_s^{\bullet,\bullet}$ are equal to impulse responses at horizon- s (Auclert et al. 2019). That is, impulse responses are sufficient to compute the elasticities algebraically, without having to simulate the model with aggregate shocks. In turn, the sequence-space Jacobian method is well-suited for rapid computation of impulse responses for different guesses of the elasticities. Implemented this way, matching business cycle moments via Simulated Method of Moments is feasible on a laptop computer, even for complex heterogeneous-agent models such as mine.

5 Quantifying the model

The quantitative model is meant to capture the extent of cyclical unemployment risk and households' exposure to it by gender and marital status in the United States. Therefore, I target the level and cyclical volatility of labor market flows in the CPS for the period between 1995-M1 and 2019-M12. In line with the flows data, I set the model frequency to one month.

5.1 Fixed parameters

The first set of parameters are calibrated externally to typical values from the literature. Table 6 summarizes my choices.

Panel A. The labor productivity process depends on gender but not marital status, because there are no recent estimates at that level of disaggregation. I normalize the average productivity of men to 1, and set the average for women to be 0.8 which is the gender wage gap observed in the CPS for 1995–2019. I set the monthly autocorrelation to be $\rho_z^m = 0.98$ for men and $\rho_z^f = 0.973$ for women, based on the estimates of Chang and Kim (2006). The standard deviation is informed by

Song, Price, Guvenen, Bloom and Von Wachter (2019), who report that the cross-sectional standard deviation of pre-tax income, in logs, is 0.943 for men, and 0.86 for women. I scale these values down by 0.82 to account for progressive taxation as estimated by Heathcote, Storesletten and Violante (2017).

Panel B. I set the real interest rate to $r = 1.92\%$ annually, and assume that households cannot borrow $\underline{a} = 0$. The correlation of separation shocks for dual-earner couples is $\rho = 0.042$, the average value in the CPS between 1995 and 2019. During the same period, 59% of individuals aged 25–54 lived in a same household with their spouse. This implies that married households account for $0.59 / (2 - 0.59) = 42\%$ of all households. The remaining mass is split equally between single men and women. I set the scale of taste shocks to $\sigma_\varepsilon(x) = 0.01$ for $x \in \{M, B\}$, and $\sigma_\varepsilon(L) = 0.04$. These are small values that induce a mild smoothing of the policy functions. Assigning larger taste shocks to UI ineligible households increase flows between unemployment and non-participation which would otherwise be minuscule in equilibrium.¹⁴

Panel C. The marginal labor tax is $\tau = 0.3$ and lump-sum transfers are negligible at $T = 0.001$. Positive transfers ensure that households can maintain a positive consumption at all times, but the precise value is not important given that I calibrate the MPCs internally. Government bonds B will be pinned down by asset demand from households. Government spending G is determined residually from the budget constraint. Unemployment benefits offer a replacement rate $b = 50\%$ up to a limit of $\bar{b} = 66\%$ of the average wage. Together with an expiration rate of $\xi = 1/6$ per month, this provides a reasonably good description of the US system in normal times.

5.2 Internally calibrated steady-state parameters

The second set of parameters are those that affect the steady state and need to be calibrated within the model. These are the type-specific utility parameters $\{\beta, \varphi, \chi\}$ and search frictions $\{f, f^*, s\}$. Given the fixed parameters, these may be calibrated separately for single men, single women, and married couples. Table 7 summarizes the results.

Identification. Although the parameters are calibrated jointly, each is identified by a particular moment in the data. I pin down the discount factor by targeting the average marginal propensity to consume (MPC). Impatient households accumulate fewer assets, which leaves them more exposed to labor income shocks. The disutility of work φ and of search χ are chosen to match labor market stocks: the mass of employed and unemployed workers. Finally, each search friction is pinned down by the transition rate it affects most directly. The job-finding rate of unemployed workers f is identified by the average UE flow, the job-finding rate of non-participants f^* by the average NE flow, and the involuntary separation rate s by the average EU flow.

Interpretation. The MPC target is 25% quarterly for all groups. Kaplan and Violante (2014) report this as the headline estimate of average MPC from a large empirical literature. The available

¹⁴In principle, one could raise taste shocks of UI ineligible households until the model matches observed NU or UN flows. However, I found that this requires very large shocks. Since these flows are subject to large measurement error, I prefer not to explicitly target them and keep taste shocks smaller.

Table 6: Fixed parameters

Parameter		Value
A. Labor productivity		
Average (M, W) ¹	\bar{z}	1, 0.8
Persistence (M, W)	ρ_z	0.98, 0.973
Cross-sectional standard deviation of log (M, W)	σ_z	0.7733, 0.7052
B. Households		
Real rate, annualized	r	1.92%
Borrowing constraint	\underline{a}	0
Correlation of spousal job loss	ρ	0.042
Population shares (SM, SW, MC) ²	ω	29%, 29%, 42%
Scale of taste shocks (M, B, L) ³	σ_ε	0.01, 0.01, 0.04
C. Government		
Labor tax	τ	0.3
Lump-sum transfer	T	0.001
UI replacement rate	b	0.5
UI cap	\bar{b}	0.66
UI expiry rate	ζ	1/6

¹ (M, W) stand for men, women.² (SM, SW, MC) stand for single men, single women, married couples.³ (M, B, L) stand for matched, UI eligible, neither.

Table 7: Internally Calibrated Steady-State Parameters

Parameter		Value	Moment	Data	Model
Single Men					
Discount factor	β	0.9659	Quarterly MPC	25%	24.99%
Disutility of work	φ	1.4352	Employment to Population	78.26%	78.29%
Disutility of search	χ	0.9029	Unemployment to Population	6.33%	6.47%
Job offer rate, U	f	0.2506	UE transition probability	25%	25.02%
Job offer rate, N	f^*	0.1208	NE transition probability	8%	8.01%
Exog separation rate	s	0.0209	EU transition probability	1.82%	1.82%
Single Women					
Discount factor	β	0.9688	Quarterly MPC	25%	24.98%
Disutility of work	φ	1.5685	Employment to Population	74.19%	74.47%
Disutility of search	χ	0.9997	Unemployment to Population	5.20%	5.57%
Job offer rate, U	f	0.2206	UE transition probability	22%	22.03%
Job offer rate, N	f^*	0.0983	NE transition probability	7%	6.89%
Exog separation rate	s	0.0150	EU transition probability	1.29%	1.29%
Married Couples					
Discount factor	β	0.9889	Quarterly MPC	25%	25.00%
Disutility of work (men)	φ	0.2157	Employment to Population	90.70%	90.73%
Disutility of search (men)	χ	0.2688	Unemployment to Population	3.04%	2.52%
Job offer rate, U (men)	f	0.2999	UE transition probability	30%	29.97%
Job offer rate, N (men)	f^*	0.1920	NE transition probability	10%	9.95%
Exog separation rate (men)	s	0.0105	EU transition probability	0.92%	0.92%
Disutility of work (women)	φ	0.3087	Employment to Population	70.38%	70.48%
Disutility of search (women)	χ	0.6602	Unemployment to Population	2.60%	1.70%
Job offer rate, U (women)	f	0.2498	UE transition probability	25%	24.97%
Job offer rate, N (women)	f^*	0.1332	NE transition probability	6%	5.76%
Exog separation rate (women)	s	0.0154	EU transition probability	0.77%	0.77%

evidence suggests that average MPC does not vary with gender and marital status. I verify this in the Italian Survey of Household Income and Wealth (SHIW).¹⁵ I find that the average MPCs of single men, single women and married couples are remarkably close, in the $[0.48, 0.50]$ range annually, despite large differences in income and assets. Note that matching the uniform MPC target requires couples to be more patient than singles. This is because spousal insurance lowers couples' need for precautionary savings. If they were as impatient as singles, they would save very little at the equilibrium interest rate, and hit the borrowing constraint often.

The labor market targets are averages from the CPS for the civilian population aged 25–54. As we saw in section 2, labor market outcomes vary strongly by gender and marital status. Notice that flows are matched almost perfectly, while the unemployment to population ratio is a little low for married workers. There is some tension between matching stocks and flows, and I put a higher weight on flows because my primary concern is labor income risk.

The importance of search frictions vary between the different flows. UE flows are essentially equal to the job-finding rate f , because workers who do not want a job are not searching in the first place. In contrast, the finding rate f^* has to be higher than the observed NE flows, because a large share of non-participants have no desire to work and reject job offers. The involuntary separation rate s is also higher than the observed EU flows. There are households whose preferences are $E \succ N \succ U$. They don't quit on their own, but exit the labor force as soon as they lose their job.

5.3 Estimated parameters

It remains to specify the elasticities of search frictions with respect to output that make unemployment risk state-dependent. I estimate these jointly by targeting an equal number of business cycle moments via Simulated Method of Moments (SMM). My novel implementation of SMM is described in section 4.2.

First, I have to take a stance on the nature of aggregate shocks that drive business cycle in the model. Although the SMM estimation itself extends trivially to multiple shocks, for simplicity I focus on a single shock that perturbs the discount factor of households. I assume that the shock follows an AR(1) process with a monthly autocorrelation of 0.95, a conventional value in light of DSGE models estimated via full information maximum likelihood methods (Justiniano, Primiceri and Tambalotti 2010, Auclert, Rognlie and Straub 2020). With a single aggregate shock process, the standard deviation of innovations is irrelevant, because I linearize the model with respect to aggregate shocks.

Conditional on the stochastic process for aggregate shocks, I can estimate the elasticities of search frictions to match the elasticities of UE , EU , and NE flows with respect to output. I choose these moments as targets because they identify the frictions most closely, and thus capture cyclical unemployment risk. As table A2 shows, the model can match all 12 moments perfectly. To the best of my knowledge, mine is the first heterogeneous-agent DSGE model that captures the cyclical

¹⁵The SHIW is an annual survey that explicitly asks correspondents about their annual MPC, which makes it well-suited to uncover heterogeneity in MPCs (Jappelli and Pistaferri 2014).

volatility of unemployment this well.

5.4 Non-targeted moments

Table 8 shows that the model does a goods job of matching moments that were not explicitly targeted in the calibration but matter for spousal insurance. The EN transition probabilities are close to the data for all types. This means that the model captures well the composition of separations, since EU transition rates were targeted and hit very closely (Table 7).

The added worker effect is defined as the increase in the probability that a married woman enters the labor force when her husband becomes unemployed. The added worker effect in my model is somewhat lower than in the data. This shows that my results are not due to exaggerating this common form of active spousal insurance.

The model captures the joint distribution of employment states in married households fairly well. It predicts too many single-earner households and thus undershoots the share of two-earner and zero-earner couples.

Table 8: Non-Targeted Moments

	Data	Model
EN transition probabilities		
Single men	1.89%	1.81%
Single women	2.06%	2.20%
Married men	0.81%	0.65%
Married women	2.59%	2.11%
Added worker effect*	7.73%	4.72%
Joint distribution of employment in married households		
(E, E)	67.19%	62.36%
(E, U)	2.01%	1.36%
(E, N)	22.29%	27.00%
(U, E)	1.90%	1.65%
(U, U)	0.26%	0.14%
(U, N)	0.66%	0.73%
(N, E)	3.47%	6.46%
(N, U)	0.18%	0.20%
(N, N)	2.14%	0.09%

Source: CPS, monthly files for 1995–2019. Sample: Civilian population aged 25–54.

* The added worker effect in the data is from [Mankart and Oikonomou \(2017\)](#) Table 5. They estimate a linear probability model on CPS data spanning the years 1994–2014 on the sample of the civilian population aged 25–55. Their sample is almost identical to what I use in this paper.

6 Macroeconomic consequences of spousal insurance

Having calibrated the steady state and estimated cyclical unemployment risk, I now use the model to isolate the role of spousal insurance in the amplification of typical business cycles with the help of a counterfactual experiment.

6.1 Household labor supply in response to idiosyncratic risk

I provide a brief overview of household labor supply in the stationary equilibrium in which there is idiosyncratic risk but no aggregate risk. This sets the stage for the discussion of spousal insurance over the business cycle in the next section.




Figure 3: Selected Labor Supply Choices of Singles

Figure 3 illustrates how singles make labor supply choices. For a matched worker, labor supply choice comes down to accepting or rejecting the offer. The left panel shows that the probability of acceptance is increasing in productivity and decreasing in assets. This is because working has a fixed utility cost while the wage is proportional to productivity. An unmatched worker can increase her labor supply by choosing unemployment over non-participation. The two panels on the right show that this decision is governed by the same logic as the acceptance of offers: poor households with high productivity are more likely to search actively. Note that the choice probability is a logistic function of assets because of the taste shocks. Otherwise, there would be a single threshold $a(z_{it})$ at which the probability jumps from 1 to 0.




Figure 4: Selected Labor Supply Choices of Couples

Figure 4 illustrates the labor supply of a typical one-earner family. For concreteness, let the husband have average productivity and the wife have slightly below average productivity. When both spouses have offers, they prefer that only the husband works, unless they are very poor. When the husband loses his job (but still collects UI), the wife's acceptance threshold increases. For assets below 2, she prefers to work until her husband finds a new job. The choices are a bit more complicated when the husband's benefits expired already. If the family is rich, they both stay out of the labor force and wait for the husband to find a job. As they are running down their assets, they become increasingly likely to send the husband to search for a job. Eventually, the wife will be willing to accept a job while the husband searches.

6.2 Household labor supply and the transmission of aggregate shocks

First, I explain why my model can generate aggregate demand-driven recessions despite having no nominal rigidities. Second, I discuss how spousal insurance affects the transmission of unemployment risk to consumption and labor supply.

Figure 5 shows the dynamic effect of a contractionary aggregate demand shock on the equilibrium interest rate and consumption by household type. The left panel shows the shock itself, a persistent rise in the discount factor of all households, which prompts them to consume less, save more, and increase their labor supply. Therefore, the ex-post interest rate has to rise on impact and then fall persistently to clear the market via income and intertemporal substitution effects. Relative to a neoclassical model, search frictions prevent supply from rising sufficiently to offset the direct effect of the shock on demand, which allows output to fall in equilibrium.¹⁶ The recession is then amplified by the ensuing increase in unemployment which depresses both z'' and supply further.

¹⁶If labor supply was a static, frictionless hours choice, households would want to raise it so much that the interest rate would have to fall enough to completely offset the effect of the discount factor shock.

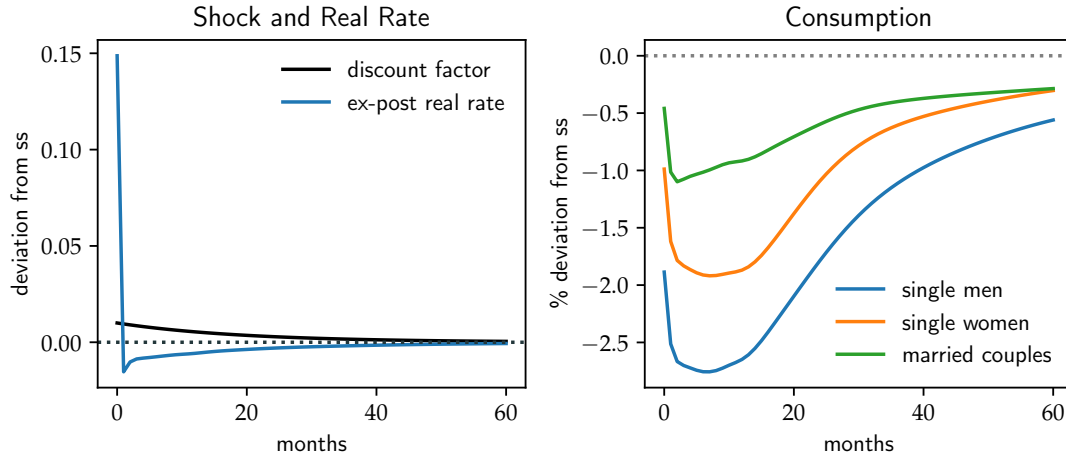


Figure 5: Impulse Responses to a Contractionary Aggregate Demand Shock

The right panel of Figure 5 shows the consumption responses broken down by household type. The large increase in financial income in period 0 mitigates but does not prevent an immediate fall in consumption, which gets worse as more workers lose their jobs and takes a long time to recover. The three types of households are affected very differently. On average, the consumption of married couples falls about half as much as the consumption of single women, and little more than third as much as the consumption of single men. Since each group has the same MPC in the stationary equilibrium, most of the difference is driven by what happens to income. Labor income is subject to both employment and wage shocks, but only the employment shocks vary with the cycle. Therefore, I turn to labor market responses next.




Figure 6: Labor Market Responses to a Contractionary Aggregate Demand Shock

Figure 6 shows how the mass of employed, unemployed, and non-participating agents evolves in response to the shock. As expected from the estimated differences in job stability, employment falls most among single men and least among married women. However, married women's employment actually increases as more women decide to enter the labor force. The same holds for

married men, albeit to a lesser extent. I argue that this is because secondary earners provide active spousal insurance by entering the labor force to compensate for job loss by the primary earner.

The steady-state labor supply policies we saw on Figures 3 and 4 are already sufficient to see why singles are more likely to drop out of the labor force in a recession, while married women (and men) may increase their labor supply. Let us start with singles. Comparing the first two panels of Figure 3 show that single men with average productivity and assets between 2.3 and 3.0 engage in job hoarding (Garibaldi and Wasmer 2005). These workers don't quit their job on their own, but will not be willing to search actively if they lose it, even if that means giving up unemployment benefits. More quit as soon as their benefits expire. In the recession, more singles lose their jobs, and the hoarders among them quit the labor force.

Figure 7 shows that the labor supply behavior of singles changes relatively little when they learn about the recession in month 0. Matched workers become marginally more likely to accept the job and less likely to refuse unemployment benefits. This is driven by intertemporal substitution: jobs are more valuable as an asset when they are harder to find. However, the same reasoning implies that ineligible singles are less likely to stay in the labor force. In sum, singles not only have more cyclical employment shocks than married men and women, but also make procyclical labor force participation decisions. This means that their income falls more in the downturn, forcing them to cut consumption more than married households.




Figure 7: Response of Singles' Labor Supply on Impact

Couples are different in that they each have a primary and a secondary earner, dictated by comparative advantage.¹⁷ Poor couples prefer to keep both of their members in the labor force, while richer couples find it optimal to keep the secondary earner out of the labor force. As Figure 4 shows, a job loss by the primary earner can induce the secondary earner to accept a job, at least temporarily. In that example, this happens for a rather large range of assets between 0.25 and 2. This explains why employment of married women can rise when more men lose their job.

¹⁷Comparative advantage depends on the labor productivity, utility cost of work, and search frictions of the spouses.




Figure 8: Response of Couples' Labor Supply on Impact

Moreover, secondary earners tend to adjust their labor supply policies in a countercyclical fashion. Figure 8 shows that when a secondary earner (in that example, the wife) gets matched in the first month of the recession, she is substantially more likely to accept regardless of the status of the primary earner. On top of the intertemporal substitution effect that I already described for singles, married households have a precautionary motive to offset the heightened risk to the primary earner.

In conclusion, married households are much less exposed to unemployment risk than singles. Spousal insurance is partly mechanical: joint income is less volatile than individual income when shocks are imperfectly correlated.¹⁸ I call this channel passive insurance, and explains why employed-employed couples (the majority of married households) are better insured than singles. However, we saw that households provide active insurance as well: secondary earners increase their labor supply to make up for realized (and anticipated) income shocks to the primary earner.

6.3 Business cycles with and without spousal insurance

We saw that the spousal insurance motive meaningfully changes the consumption and labor supply of married households relative to singles. But how much does this matter for aggregate outcomes? In this section, I answer this question with the help of a counterfactual experiment.

No spousal insurance. To isolate the role of spousal insurance in the contemporary US economy, I consider a counterfactual economy where all married households have split up. The formerly married men and women have the same income risk and MPC as before, which maximizes comparability to the benchmark model. The only parameters I recalibrate are the discount factor β , disutility of work φ , and disutility of search χ . These are necessary for the separated men and women to have the same MPC, employment rate, and unemployment rate as before. That is, the counterfactual economy is as close as possible at both the individual and macro levels, but differs in terms of household formation.

¹⁸If both wage and employment shocks were perfectly correlated across spouses, a married household would simply be a scaled-up version of a single household.




Figure 9: Consumption Responses to the Same Aggregate Shock

Figure 9 plots the consumption responses to the same aggregate demand shock in the counterfactual model without spousal insurance against the benchmark model. Starting with the right panel, the consumption of married men and women fall much more without spousal insurance, with a peak response that is almost 3 times as large in percentage terms as in the benchmark case. Unsurprisingly, taking away an insurance mechanism from these households leaves them more exposed to the recession. However, single men and women also suffer larger drops in consumption in the economy without spousal insurance. The reason is that when demand from married men and women falls, it aggravates unemployment for every other household in the economy. My model captures this spillover accurately via the estimated gender and marital status-specific search frictions. Spousal insurance weakens the feedback between unemployment and aggregate demand, and thereby stabilizes the macroeconomy more broadly.

7 Spousal insurance across time

The stabilizing potential of spousal insurance can vary substantially over time. In this section, I use my model to characterize the effect of three scenarios on spousal insurance. The first two are long-run trends: the decline in marriage rate, and the rise in female labor force participation. The third is the incidence of job losses during the early months of the COVID-19 recession.

7.1 Marriage rate

Household formation determines how widespread the access to spousal insurance is. The left panel of Figure 10 shows that the share of married households in the United States has fallen from about 70% in the 1960s to less than 40% in 2019. This is a dramatic change in the composition of households.

Interpreting the decline in marriage rate is subject to the caveat that it overstates the decline in spousal insurance because of the secular rise in cohabitation, i.e. the number of unmarried people living together. Unfortunately, cohabitation has been measured consistently in the CPS since 1994

only. Between 1994 and 2018, the share of cohabiting adults aged 25–54 increased from 3.6% to 9.3%, which is not nearly enough to offset the decline in marriage, especially compared to its high level in 1960s.

To quantify the impact on macroeconomic volatility, I change the household composition in the model to represent 1960s. The shaded areas represent the basis for the benchmark calibration (blue) and counterfactual calibration (green). Then, I subject the model to the same aggregate demand shock as before. The right panel of Figure 10 shows that a larger share of married households, all else equal, leads to a much smaller and less persistent fall in aggregate consumption.

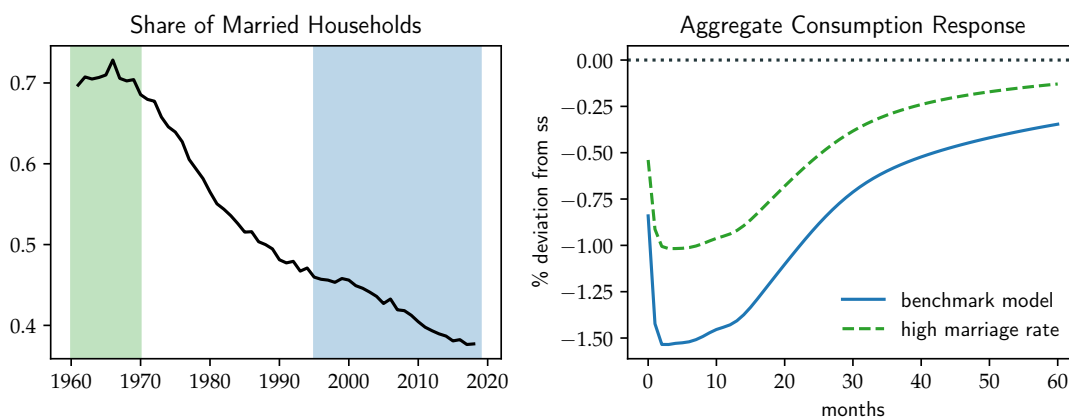


Figure 10: Marriage Rate and Aggregate Volatility

Share of married households is from the CPS ASEC. Sample is restricted to civilian population aged 25-54.

Note that this experiment not only extends the availability of spousal insurance but lowers the amount of individual unemployment risk in the economy as well. This is because married couples have lower probability of job loss than singles. My modeling choice is justified by the time series of *EU* and *UE* flows which have not diverged between married and single workers despite the steep and steady increase in the share of single households (Figure 2). But it is also a conservative choice, since giving spousal insurance would be even more valuable to people with higher individual risk.

7.2 Female Labor Force Participation

The left panel of Figure 11 shows that the labor force participation rate of married women in the United States increased from less than 40% in 1960 to more than 70% by 1995, where it has plateaued. This is a profound structural change that has a subtle impact on spousal insurance.

On the one hand, women have had more stable jobs and more countercyclical labor supply than men throughout this period. Therefore, one may expect that aggregate volatility will decrease as more for them decide to work. [Albanesi \(2019\)](#) makes this argument for why women's employment contributed to the Great Moderation. On the other hand, when more married women

are already in the labor force, the room for countercyclical entry diminishes. The idea that dual-earner families may be more vulnerable than single-earner families is emphasized by [Warren and Warren Tyagi \(2004\)](#). My model accommodates both of these considerations, and therefore can provide a nuanced answer to this question.




Figure 11: Female Labor Force Participation and Aggregate Volatility

Labor force participation rates are from the CPS ASEC. Sample is restricted to civilian population aged 25-54.

I calibrate the economy to the 1960s level of female labor force participation with a combination of lower wages and higher disutility of work for women. The gender wage gap is an observed moment. [Doepke and Tertilt \(2016\)](#) report that the gap in the median wage was 45% in 1960. I find that targeting this higher gender wage gap in the model explains about half the rise in married women's participation. I use the disutility of employment to make up for the rest.¹⁹ Finally, I recalibrate the discount factor to keep the quarterly MPC at 25% and the disutility of search to target an unemployment rate of 3.5% for married women.

The right panel of Figure 11 shows that low female labor force participation (FLFP) has an ambiguous effect on the volatility of aggregate consumption. In the first three months, consumption falls less in the benchmark model, which means that high FLFP does mitigate the immediate impact of the shock. However, the subsequent recovery is lower in the modern economy. This pattern reflects the offsetting effects of FLFP on the passive and active channels of spousal insurance. Passive insurance acts without delay and benefits only dual earner couples. There are more of them in the modern economy. Active insurance is more relevant for single-earner families, whose secondary earner is out of the labor force. Her entry is slowed down by the search frictions, but only for a few months. In sum, the model implies that high female labor force participation boosts passive insurance at the cost of weakening active insurance, and ultimately has a small impact on aggregate volatility.

¹⁹Disutility of work may be thought of as capturing various factors behind the rise of female labor force participation during this period such as the increased availability of oral contraception ([Goldin and Katz 2002](#)) and household appliances ([Greenwood, Seshadri and Yorukoglu 2005](#)) as well as cultural change ([Fernández 2013](#)).

Takeaway. My model accounts for the fact that women have less procyclical employment than men. The difference is partly targeted (estimated elasticities of separation) and partly micro-founded (active spousal insurance). The experiment demonstrates that these differences are not sufficient to conclude that high FLFP dampens business cycle fluctuations.

My model highlights a key moment that drives the net effect of FLFP on macroeconomic volatility. It is how easy it is for married women out of the labor force to find a job when they want one. If it was impossible, FLFP would be synonymous with spousal insurance, because one-earner couples would behave as singles. If labor supply was frictionless, FLFP would be irrelevant for spousal insurance, since a non-participant spouse could earn income at will. My model is between these two extremes. In particular, I keep search frictions constant at a moderate level. This fits well the evolution of gross flows which shows that declining NE transitions are the dominant driver of rising participation of married women (Figure 2).

This experiment, however, is not a conclusive proof that the rise in FLFP did not have a dampening effect. A caveat to my analysis is that married women today may have a much higher earnings capacity even when they are out of the labor force, as a result from extra work experience they accumulate before quitting the labor force. Although my model does have the property that people with low productivity select to stay out of the labor force, it does not include explicit skill depreciation.

7.3 COVID-19 recession

Spousal insurance is particularly powerful against unemployment because of two regularities of cyclical job loss. First, the correlation of job loss between spouses is essentially zero. Second, the probability of job loss is much less cyclical for women than men.

Figure 12 shows that job losses during the early months of COVID-19 pandemic did not follow the usual pattern.²⁰ First, married women were much more likely to lose their job than married men. Second, the correlation of spousal job loss was unusually high during this period. In fact, both EU flows and the correlation of spousal job loss reached their all time high in April 2020, at least in the four decades covered by the CPS sample.

²⁰ Alon et al. (2020a) were the first to point out that COVID-19 recession hit women particularly hard.




Figure 12: Unemployment Risk of Married Households in 2020

Source: CPS monthly files. Sample is restricted to civilian population aged 25-54.

To assess the effect of this large and sudden shock on household consumption, I pursue a different strategy than before. Instead of generating a recession via a reduced-form discount factor shock, I feed the observed paths of EU flows, UE flows, and correlation directly into the decision problem of married households. Specifically, I assume that the economy starts from its stationary equilibrium in February 2020. In March, households are surprised by the low job-finding rates and high separation rates, but believe that these are temporary deviations and April will be normal. In April, they are surprised again. The same process continues until August 2020.

Figure 13 shows the average consumption response (red line). To put this response in context, I construct a “regular recession” as follows. Married men face the same transition probabilities as before. However, the probability of married women’s job loss follows the usual pattern, and rises just $7.87/13.6 = 58\%$ as much as married men’s. In addition, the correlation of joint job loss remains constant at 0.042. Although this counterfactual recession is still very severe, household consumption falls less than three-quarters as much than in the COVID-19 recession.




Figure 13: Consumption Response to Early Months of COVID-19 Recession

Takeaway. The example of the COVID-19 recession shows that the nature, and not just the size,

of an economic shock matters for its impact on households. Aggregate shocks that originate in sectors with a high employment share of just one gender (such as construction) are better insured than shocks that affect women and men equally and are more likely to affect two spouses at the same time.

8 Conclusion

This paper was motivated by the observation that married households rely heavily on spousal labor supply to absorb idiosyncratic income shocks. I built a general equilibrium model with incomplete markets in which households have access to this important channel of partial insurance. I estimated the model on US data, so that it reproduced the volatility of cyclical unemployment risk and labor force participation patterns by gender and marital status. Then, I used the model to argue that spousal insurance dampens aggregate fluctuations substantially via both demand and supply channels.

An immediate consequence of this result is that business cycle analysis ought to be broader in scope. Any feature of the economy that boosts spousal labor supply can be expected to dampen business cycles. Conversely, anything that hinders spousal insurance has the potential to amplify business cycles. I called attention to long-term trends as well as aggregate shocks that have this property.

Future research could expand this project in two directions. First, a case could be made for rethinking social policies that provide more insurance to married households than singles. Married couples have more choices than singles to maintain health insurance coverage, to collect social security benefits, and to optimize saving for retirement. Thus, the tax and transfer system reinforces the insurance advantage that married couples already have from spousal labor supply. The design of these policies may be rooted in the old model of single earner families, or chosen with long-term goals in mind, such as promoting marriage. My results suggest that business cycle stabilization calls for using the tax and transfer system to provide extra insurance to single households instead.

Second, my model of the household is stylized in many dimensions. The family economics literature has emphasized that assortative mating with respect to education (long-term earnings potential) is an important driver of inequality (Fernandez, Guner and Knowles 2005). Assortative mating is likely to play a role in the business cycle context as well. Here, the relevant notion is assortative mating with respect to job stability and short-run earnings potential, i.e. the determinants of spousal insurance. In addition, I modeled households as permanent units that make decisions jointly. But limits of cooperative behavior could have first-order effects on family labor supply and consumption smoothing, and hence on the macroeconomy as well. I leave the development of richer models along these lines for future research.

References

- Abowd, John M and Arnold Zellner**, “Estimating Gross Labor-Force Flows,” *Journal of Business & Economic Statistics*, 1985, 3 (3), 254–283.
- Ahn, SeHyouon, Greg Kaplan, Benjamin Moll, Thomas Winberry, and Christian Wolf**, “When Inequality Matters for Macro and Macro Matters for Inequality,” in “NBER Macroeconomics Annual 2018, Volume 32,” University of Chicago Press, 2018.
- Albanesi, Stefania**, “Changing Business Cycles: The Role of Women’s Employment,” Working Paper 25655, National Bureau of Economic Research 2019.
- **and Ayşegül Şahin**, “The Gender Unemployment Gap,” *Review of Economic Dynamics*, 2018, 30, 47–67.
- Alon, Titan, Matthias Doepke, Jane Olmstead-Rumsey, and Michèle Tertilt**, “The Impact of COVID-19 on Gender Equality,” *Covid Economics: Vetted and Real-Time Papers*, 2020.
- , — , — , **and** — , “This Time It’s Different: The Role of Women’s Employment in a Pandemic Recession,” Working Paper 27660, National Bureau of Economic Research 2020.
- Alves, Felipe**, “Job Ladder and Business Cycles,” Job Market Paper, NYU 2019.
- Auclert, Adrien, Bence Bardóczy, Matthew Rognlie, and Ludwig Straub**, “Using the Sequence-Space Jacobian to Solve and Estimate Heterogeneous-Agent Models,” Working Paper 26123, National Bureau of Economic Research 2019.
- , **Matthew Rognlie, and Ludwig Straub**, “The Intertemporal Keynesian Cross,” Working Paper 25020, National Bureau of Economic Research 2018.
- , — , **and** — , “Micro Jumps, Macro Humps: Monetary Policy and Business Cycles in an Estimated HANK Model,” Working Paper 26647, National Bureau of Economic Research 2020.
- Bayer, Christian and Ralph Luetticke**, “Solving Heterogeneous Agent Models in Discrete Time with many Idiosyncratic States by Perturbation Methods,” *Quantitative Economics*, 2020, *forthcoming*.
- , **Ralph Lütticke, Lien Pham-Dao, and Volker Tjaden**, “Precautionary Savings, Illiquid Assets, and the Aggregate Consequences of Shocks to Household Income Risk,” *Econometrica*, 2019, 87 (1), 255–290.
- Birinci, Serdar**, “Spousal Labor Supply Response to Job Displacement and Implications for Optimal Transfers,” Working Paper, Federal Reserve Bank of St. Louis 2019.
- Blundell, Richard, Luigi Pistaferri, and Ian Preston**, “Consumption Inequality and Partial Insurance,” *American Economic Review*, 2008, 98 (5), 1887–1921.

- , —, and Itay Saporta-Eksten, “Consumption Inequality and Family Labor Supply,” *American Economic Review*, 2016, 106 (2), 387–435.
- Boppart, Timo, Per Krusell, and Kurt Mitman**, “Exploiting MIT Shocks in Heterogeneous-Agent Economies: The Impulse Response as a Numerical Derivative,” *Journal of Economic Dynamics and Control*, 2018, 89, 68–92.
- Busch, Christopher, David Domeij, Fatih Guvenen, and Rocio Madera**, “Skewed Idiosyncratic Income Risk over the Business Cycle: Sources and Insurance,” Manuscript 2020.
- Carroll, Christopher D**, “The Method of Endogenous Gridpoints for Solving Dynamic Stochastic Optimization Problems,” *Economics letters*, 2006, 91 (3), 312–320.
- Challe, Edouard, Julien Matheron, Xavier Ragot, and Juan F. Rubio-Ramirez**, “Precautionary Saving and Aggregate Demand,” *Quantitative Economics*, 2017, 8 (2), 435–478.
- Chang, Yongsung and Sun-Bin Kim**, “From Individual to Aggregate Labor Supply: A Quantitative Analysis Based on a Heterogeneous Agent Macroeconomy,” *International Economic Review*, 2006, 47 (1), 1–27.
- Doepke, Matthias and Michèle Tertilt**, “Families in Macroeconomics,” in “Handbook of Macroeconomics,” Vol. 2, Elsevier, 2016, pp. 1789–1891.
- Druedahl, Jeppe**, “A Guide On Solving Non-Convex Consumption-Saving Models,” Manuscript 2020.
- Ellieroth, Kathrin**, “Spousal Insurance, Precautionary Labor Supply, and the Business Cycle,” Job Market Paper, Indiana University 2019.
- Elsby, Michael WL, Bart Hobijn, and Ayşegül Şahin**, “On the Importance of the Participation Margin for Labor Market Fluctuations,” *Journal of Monetary Economics*, 2015, 72, 64–82.
- Fernández, Raquel**, “Cultural Change as Learning: The Evolution of Female Labor Force Participation over a Century,” *American Economic Review*, 2013, 103 (1), 472–500.
- Fernandez, Raquel, Nezih Guner, and John Knowles**, “Love and Money: A Theoretical and Empirical Analysis of Household Sorting and Inequality,” *The Quarterly Journal of Economics*, 2005, 120 (1), 273–344.
- Ferra, Sergio De, Kurt Mitman, and Federica Romei**, “Household Heterogeneity and the Transmission of Foreign Shocks,” *Journal of International Economics*, 2020, p. 103303.
- Flood, Sarah, Miriam King, Renae Rodgers, Steven Ruggles, and Robert J. Warren**, “Integrated Public Use Microdata Series, Current Population Survey: Version 7.0 [dataset],” Minneapolis, MN: IPUMS. <https://doi.org/10.18128/D030.V7.0> 2020.

- Fukui, Masao, Emi Nakamura, and Jón Steinsson**, “Women, Wealth Effects, and Slow Recoveries,” Working Paper 25311, National Bureau of Economic Research 2018.
- Garibaldi, Pietro and Etienne Wasmer**, “Equilibrium Search Unemployment, Endogenous Participation, and Labor Market Flows,” *Journal of the European Economic Association*, 2005, 3 (4), 851–882.
- Goldin, Claudia and Lawrence F Katz**, “The Power of the Pill: Oral Contraceptives and Women’s Career and Marriage Decisions,” *Journal of Political Economy*, 2002, 110 (4), 730–770.
- Gorbachev, Olga**, “Has the Increased Attachment of Women to the Labor Market Changed a Family’s Ability to Smooth Income Shocks?,” *American Economic Review*, 2016, 106 (5), 247–51.
- Gornemann, Nils, Keith Kuester, and Makoto Nakajima**, “Doves for the Rich, Hawks for the Poor? Distributional Consequences of Monetary Policy,” Working Paper, Federal Reserve Bank of Philadelphia 2016.
- Graves, Sebastian**, “Does Unemployment Risk Affect Business Cycle Dynamics?,” Job Market Paper, NYU 2019.
- Greenwood, Jeremy, Ananth Seshadri, and Mehmet Yorukoglu**, “Engines of Liberation,” *The Review of Economic Studies*, 2005, 72 (1), 109–133.
- Guerrieri, Veronica and Guido Lorenzoni**, “Credit Crises, Precautionary Savings, and the Liquidity Trap,” *Quarterly Journal of Economics*, 2017, 132 (3), 1427–1467.
- Guler, Bulent, Fatih Guvenen, and Giovanni L Violante**, “Joint-Search Theory: New Opportunities and New Frictions,” *Journal of Monetary Economics*, 2012, 59 (4), 352–369.
- Guvenen, Fatih, Fatih Karahan, Serdar Ozkan, and Jae Song**, “What Do Data on Millions of U.S. Workers Reveal about Life-Cycle Earnings Dynamics?,” Working Paper 20913, National Bureau of Economic Research 2015.
- , **Serdar Ozkan, and Jae Song**, “The Nature of Countercyclical Income Risk,” *Journal of Political Economy*, 2014, 122 (3), 621–660.
- Haan, Wouter J Den, Pontus Rendahl, and Markus Riegler**, “Unemployment (Fears) and Deflationary Spirals,” *Journal of the European Economic Association*, 2017, 16 (5), 1281–1349.
- Hall, Robert E**, “High Discounts and High Unemployment,” *American Economic Review*, 2017, 107 (2), 305–30.
- Heathcote, Jonathan, Kjetil Storesletten, and Giovanni L Violante**, “Consumption and Labor Supply with Partial Insurance: An Analytical Framework,” *American Economic Review*, 2014, 104 (7), 2075–2126.

- , —, and —, “Optimal Tax Progressivity: An Analytical Framework,” *The Quarterly Journal of Economics*, 2017, 132 (4), 1693–1754.
- Hoffmann, Eran B and Davide Malacrino**, “Employment Time and the Cyclicalities of Earnings Growth,” *Journal of Public Economics*, 2019, 169, 160–171.
- Iskhakov, Fedor, Thomas H Jørgensen, John Rust, and Bertel Schjerning**, “The Endogenous Grid Method for Discrete-Continuous Dynamic Choice Models with (or without) Taste Shocks,” *Quantitative Economics*, 2017, 8 (2), 317–365.
- Jappelli, Tullio and Luigi Pistaferri**, “Fiscal Policy and MPC Heterogeneity,” *American Economic Journal: Macroeconomics*, 2014, 6 (4), 107–36.
- Justiniano, Alejandro, Giorgio E Primiceri, and Andrea Tambalotti**, “Investment Shocks and Business Cycles,” *Journal of Monetary Economics*, 2010, 57 (2), 132–145.
- Kaplan, Greg and Giovanni L Violante**, “A Model of the Consumption Response to Fiscal Stimulus Payments,” *Econometrica*, 2014, 82 (4), 1199–1239.
- , **Benjamin Moll, and Giovanni L Violante**, “Monetary Policy According to HANK,” *American Economic Review*, 2018, 108 (3), 697–743.
- Kekre, Rohan**, “Unemployment Insurance in Macroeconomic Stabilization,” Manuscript 2019.
- Keynes, John Maynard**, *The General Theory of Employment, Interest, and Money*, Springer, 1936.
- Krusell, Per and Anthony A Smith Jr**, “Income and Wealth Heterogeneity in the Macroeconomy,” *Journal of Political Economy*, 1998, 106 (5), 867–896.
- Lundberg, Shelly**, “The Added Worker Effect,” *Journal of Labor Economics*, 1985, 3 (1, Part 1), 11–37.
- Mankart, Jochen and Rigas Oikonomou**, “The Rise of the Added Worker Effect,” *Economics Letters*, 2016, 143, 48–51.
- and —, “Household Search and the Aggregate Labour Market,” *The Review of Economic Studies*, 2017, 84 (4), 1735–1788.
- McKay, Alisdair and Ricardo Reis**, “The Role of Automatic Stabilizers in the US Business Cycle,” *Econometrica*, 2016, 84 (1), 141–194.
- Oh, Hyunseung and Ricardo Reis**, “Targeted Transfers and the Fiscal Response to the Great Recession,” *Journal of Monetary Economics*, 2012, 59, S50–S64.
- Ortigueira, Salvador and Nawid Siassi**, “How Important is Intra-Household Risk Sharing for Savings and Labor Supply?,” *Journal of Monetary Economics*, 2013, 60 (6), 650–666.

- Patterson, Christina**, “The Matching Multiplier and the Amplification of Recessions,” Job Market Paper, MIT 2018.
- Pruitt, Seth and Nicholas Turner**, “Earnings Risk in the Household: Evidence from Millions of US Tax Returns,” *American Economic Review: Insights*, June 2020, 2 (2), 237–54.
- Ravn, Morten O and Vincent Sterk**, “Job Uncertainty and Deep Recessions,” *Journal of Monetary Economics*, 2017, 90, 125–141.
- Reiter, Michael**, “Recursive Computation of Heterogeneous Agent Models,” Mimeo, Universitat Pompeu Fabra 2002.
- , “Solving Heterogeneous-Agent Models by Projection and Perturbation,” *Journal of Economic Dynamics and Control*, 2009, 33 (3), 649–665.
- Shore, Stephen H and Todd Sinai**, “Commitment, Risk, and Consumption: Do Birds of a Feather Have Bigger Nests?,” *The Review of Economics and Statistics*, 2010, 92 (2), 408–424.
- Song, Jae, David J Price, Fatih Guvenen, Nicholas Bloom, and Till Von Wachter**, “Firming up Inequality,” *The Quarterly Journal of Economics*, 2019, 134 (1), 1–50.
- Warren, Elizabeth and Amelia Warren Tyagi**, *The Two Income Trap: Why Middle-Class Parents are Going Broke*, Basic Books, 2004.
- Winberry, Thomas**, “A Method for Solving and Estimating Heterogeneous Agent Macro Models,” *Quantitative Economics*, 2018, 9 (3), 1123–1151.

A Additional Tables

Table A1: Cyclicalities of Labor Market Stocks

		single men	single women	married men	married women
Average	<i>E</i>	78.43%	74.31%	90.68%	70.37%
	<i>U</i>	6.28%	5.18%	3.04%	2.60%
	<i>N</i>	15.28%	20.52%	6.28%	27.04%
Cyclical	<i>E</i>	1.44*** (0.22)	0.82*** (0.18)	0.84** (0.12)	0.39*** (0.09)
	<i>U</i>	-1.18*** (0.17)	-0.66*** (0.09)	-0.75*** (0.11)	-0.40*** (0.05)
	<i>N</i>	-0.27 (0.09)	-0.16 (0.12)	-0.10 (0.45)	0.00 (0.08)

Source: Current Population Survey, monthly files for 1995–2019.

Sample: Civilian population aged 25–54.

Cyclical: semi-elasticity to cyclical component of real GDP. HAC standard errors in parentheses.

***: p -value < 0.01, **: p -value < 0.05.

This table summarizes the cyclical behavior of employment, unemployment, and non-participation, disaggregated by gender and marital status. Notably, non-participation is less procyclical for married people than single people (also for women than for men). The model rationalizes this outcome as a result countercyclical entry by secondary earners.

Table A2: Estimated Elasticities of Transition Probabilities

Parameter		Value	Moment	Data	Model
Single Men					
Job offer rate, U	$\varepsilon(f)$	10.69	UE flow	10.71	10.71
Job offer rate, N	$\varepsilon(f^*)$	5.57	NE flow	4.74	4.74
Separation rate	$\varepsilon(s)$	-11.17	EU flow	-10.12	-10.12
Single Women					
Job offer rate, U	$\varepsilon(f)$	8.72	UE flow	8.72	8.72
Job offer rate, N	$\varepsilon(f^*)$	3.55	NE flow	3.74	3.74
Separation rate	$\varepsilon(s)$	-6.55	EU flow	-5.68	-5.68
Married Men					
Job offer rate, U	$\varepsilon(f)$	10.37	UE flow	10.36	10.36
Job offer rate, N	$\varepsilon(f^*)$	9.74	NE flow	2.78	2.78
Separation rate	$\varepsilon(s)$	-13.35	EU flow	-13.60	-13.60
Married Women					
Job offer rate, U	$\varepsilon(f)$	8.40	UE flow	8.40	8.40
Job offer rate, N	$\varepsilon(f^*)$	10.72	NE flow	2.30	2.30
Separation rate	$\varepsilon(s)$	-3.71	EU flow	-7.87	-7.87

This table summarizes the estimated elasticities of job-finding and separation rates to output. The elasticities are estimated jointly in an exactly-identified Simulated Method of Moments, and are matched perfectly. Business cycle fluctuations in the model originate in discount factor shocks.

B Data

In this appendix, I describe the two datasets I use throughout the paper. The aggregated earnings data is published by [Pruitt and Turner \(2020\)](#) and can be downloaded from the AEA's website.²¹ The harmonized CPS micro data files can be downloaded from IPUMS.²²

B.1 Administrative earnings data

[Pruitt and Turner \(2020\)](#) start from population-level data on earnings and marriage from the Internal Revenue Service (IRS) combined with demographic information from the Social Security Administration (SSA). The gross earnings data is from the W-2 forms that include wage and salary earnings. The earnings reported on these forms are verified by employers and are not top-coded. Missing data are likely to represent true \$0 earnings.

²¹<https://www.aeaweb.org/articles?id=10.1257/aeri.20190096>

²²<https://cps.ipums.org/cps/>

To construct their sample, they draw a 1-in-5 random sample of males aged between 25 and 60. They pull the W-2s of these males and their spouses for the years 1999–2014. Then, they impose the following restrictions.

1. When an individual dies in the sample period, drop observations from the last two years of life. (remove large health-induced earnings shocks)
2. Remove individuals who receive Social Security disability payments. (focus on labor market earnings, not transfers)
3. Remove individuals with nonzero self-employment income (self-employment income is not third-party verified).
4. Require spouses to file jointly in all years used to measure earnings growth. (eliminate the effect of divorce)

This procedure leaves about 90% of joint tax returns, and more than 235 million person-year observations. Nominal earnings are converted to 2014 dollars using the CPI.

In section 2, I present selected moments of annual labor earnings growth. These can be found in file DELTA1.xlsx. Let y_{it} denote the gross labor earnings of individual i in year t , where \$0 earnings are replaced by \$1. Annual labor earnings growth is then defined as $x_{it} \equiv \log y_{it} - \log y_{it-1}$. Household-level earnings growth is given by $x_{ht} = \log(y_{ht}^m + y_{ht}^f) - \log(y_{ht-1}^m + y_{ht-1}^f)$, where y_{ht}^m and y_{ht}^f are the earnings of the male and the female spouse within the household.²³

B.2 Current Population Survey

The CPS is the primary source of labor force statistics for the population of the United States. It is conducted at the household level, and provides information about all household members. I use the basic monthly files for the years 1976–2020 and the Annual Social and Economic Supplement (ASEC) for the years 1962–2019. The monthly files are the source of worker flows. The ASEC provides a longer time-series of labor force participation rates and wages. In all cases, I restrict the sample to the civilian population aged 25–54.

The basic monthly files have a rotating panel structure. Households are interviewed for four months, have eight months off, then are interviewed again for four months. Therefore, in principle, 75% of households are observed in two consecutive months. I link individuals based on their unique identifier. I validate matches using gender, age, and marital status. Keeping matches where marital status remains constant is in line with my theoretical model, in which it is a permanent type.

Given the linked sample, the transition probability between, say, employment and unemployment in month t can be computed as follows. Count people who report being employed in $t - 1$

²³Same-sex marriage has been legal in all 50 states since 2015. The sample predates this year, hence does not include same-sex couples.

and unemployed in t , then divide by the number of all people who report being employed in $t - 1$. By construction, the number is between 0 and 1. To make the flows representative of the intended population (single men etc.), I weight by the time- t sample weights.

When estimating the elasticities to GDP, I take quarterly averages of the monthly flows, and seasonally adjust them using the X-13ARIMA-SEATS program from the Census Bureau (using the default parameters).

C Derivations

C.1 Parameterizing job loss for married couples

Let X and Y be two Bernoulli random variables given by

$$\Pr(X = 1, Y = 1) = a \quad (18)$$

$$\Pr(X = 1, Y = 0) = b \quad (19)$$

$$\Pr(X = 0, Y = 1) = c \quad (20)$$

$$\Pr(X = 0, Y = 0) = d \quad (21)$$

The total probability that either event happens is

$$\Pr(X = 1) = a + b \equiv p \quad (22)$$

$$\Pr(Y = 1) = a + c \equiv q \quad (23)$$

and their correlation is

$$\text{Corr}(X, Y) = \frac{a - pq}{\sqrt{p(1-q)q(1-p)}}. \quad (24)$$

This means that the transitions of an employed-employed couple can be constructed from the separation rate of married men (s_m), the separation rate of married women (s_f), and the correlation of job loss between spouses (ρ) as follows:

$$\Pr(\vec{EB}_m, \vec{EB}_f) = s_m s_f + \rho \sqrt{s_m s_f (1 - s_m)(1 - s_f)} \quad (25)$$

$$\Pr(\vec{EB}_m, \vec{EM}_f) = s_m (1 - s_f) - \rho \sqrt{s_m s_f (1 - s_m)(1 - s_f)} \quad (26)$$

$$\Pr(\vec{EM}_m, \vec{EB}_f) = (1 - s_m) s_f - \rho \sqrt{s_m s_f (1 - s_m)(1 - s_f)} \quad (27)$$

$$\Pr(\vec{EM}_m, \vec{EM}_f) = (1 - s_m)(1 - s_f) + \rho \sqrt{s_m s_f (1 - s_m)(1 - s_f)} \quad (28)$$